Article Recommender System

DS300.Q11 - Final Project

Presented by:
T. Nhat[1] L. M. Nhut[1] N. H. Phat[1]
Instructor: Master. Huynh Van Tin[1]

[1]Department of Computer Science
University of Information Technology

DS300 Final Project Report

Table of Contents

Table of Contents

Motivation: Why News Recommendation?

In the digital era, news platforms face **Information Overload**:

- **User Experience:** Navigating thousands of daily articles is overwhelming.
- **Retention:** Personalized content increases user engagement and platform stickiness.
- **Monetization:** Better engagement leads to higher ad revenue and subscriptions.

Target: Build a system that understands both *Content semantics* and *User behavior*.

Research Objectives

Our goal is to develop a state-of-the-art recommendation engine for Vietnamese news:

- **Hybrid Logic:** Combine Collaborative Filtering (CF) with Content-Based (CB) techniques.
- **Graph Learning:** Leverage GNNs to model complex user-article interactions.
- **Multimodal Fusion:** Incorporate Social (replies) signals.
- **Vietnamese NLP:** Optimize PhoBERT embeddings for news semantics.

Problem Statement

Challenge: How to recommend relevant news articles to users?

Key Challenges:

1. **Cold Start:** New users/articles have no interaction history
2. **Sparse Data:** Most users read very few articles (99.9% sparsity)
3. **Vietnamese Language:** Limited NLP resources for Vietnamese text
4. **Dynamic Content:** News articles are time-sensitive

Our Approach: Graph-based Contrastive Learning with PhoBERT embeddings

Table of Contents

VnExpress Data Crawling

Data Source: VnExpress.net (largest Vietnamese news portal)

Categories Crawled:

- 'Thế giới'
- 'Giáo dục'
- 'Thể thao'
- 'Thời sự'
- 'Kinh doanh'
- 'Khoa học công nghệ'

Collected Data:

- Articles: title, content, author, publish date
- Metadata: category, tags (per article)
- Comments: user ID, comment text, reactions, replies
- User profiles: user ID, username, join date

Data Collection Pipeline: 3-Step Workflow

Our crawler implements a resilient 3-stage architecture to handle high-volume news data:

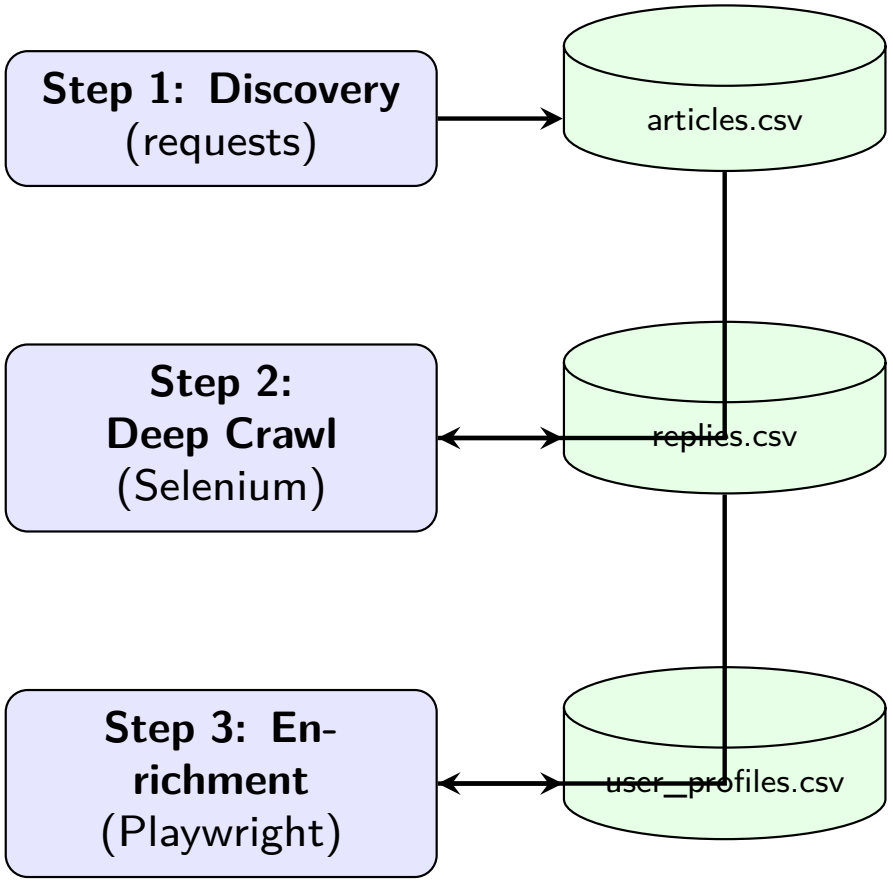

Parallel scanning of category pages to find new URLs.

Parallel browser workers to extract deep comment threads & nested replies.

Scraping full history for all unique active commenters.

Pipeline Orchestration & Efficiency

- **Parallelism:** `Step 2` (Selenium) is resource-intensive. We split the workload across N parallel browser instances using `multiprocessing`.
- **Resiliency:**
  - **Caching:** Avoids redundant HTTP requests for static content.
  - **Checkpointing:** Atomic writes to CSV allow resuming from the last article in case of failure.
  - **Headless Execution:** Minimizes resource overhead in production environments.
- **Tool Stack:**
  - **Requests:** High-speed metadata extraction.
  - **Selenium:** Handling complex JS (Load More buttons).
  - **Playwright:** Modern, fast engine for user history scraping.

Dataset Overview: High-Level Statistics

Raw Behavioral Data:

- **3,685** Articles
- **42,050** Comments
- **6** News Categories
- **19,842** User Profiles

Graph Interaction Set (G1-G3):

- **13,117** Unique Users
- **3,645** Active Articles
- **38,505** Total Interactions
- **Sparsity: 99.919%**

Dataset Overview (VnExpress)

Statistics:

- **Users:** 13,117
- **Articles:** 3,645
- **Interactions:** 38,505
- **Sparsity:** 99.92%

The dataset is highly sparse, typical of news domains where users read only a tiny fraction of daily content.

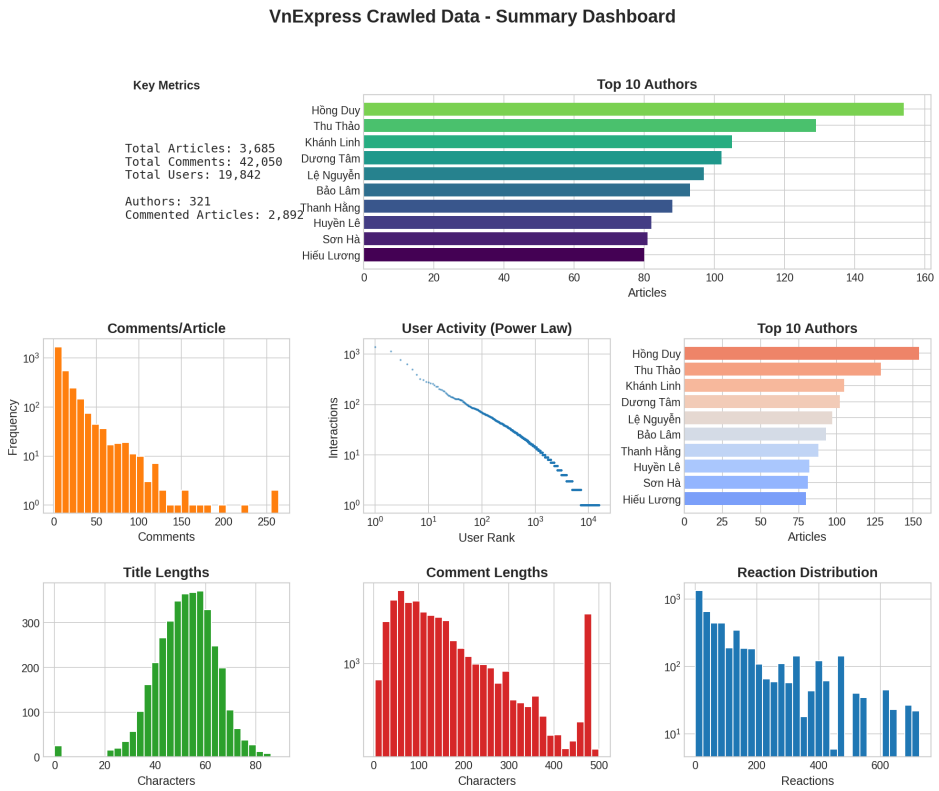

Data Sample: Interaction & Content

1. **Content Sample (Article)**
"Chinh phục loạt 'ông lớn' công nghệ sau hơn 500 lần bị từ chối"
→ **Tags:** [Meta, Amazon, Apple, Tesla, kĩ sư công nghệ]

2. **Social Sample (Comment/Reply)**
   - **User A:** "Nếu giỏi thì cố gắng mua lại mấy công ty này..."
   - **User B:** (Reply to A) "Mua lại Meta chắc cần nghìn tỷ USD..."
   ⇒ Social Connection detected

Feature Mapping

- **Article** → PhoBERT Embedding ($d = 64$) → Node Feature.
- **User** → Interaction aggregation → Node Embedding.
- **Reply** → User-User Edge in the Heterogeneous Graph.

Category Distribution (VnExpress)

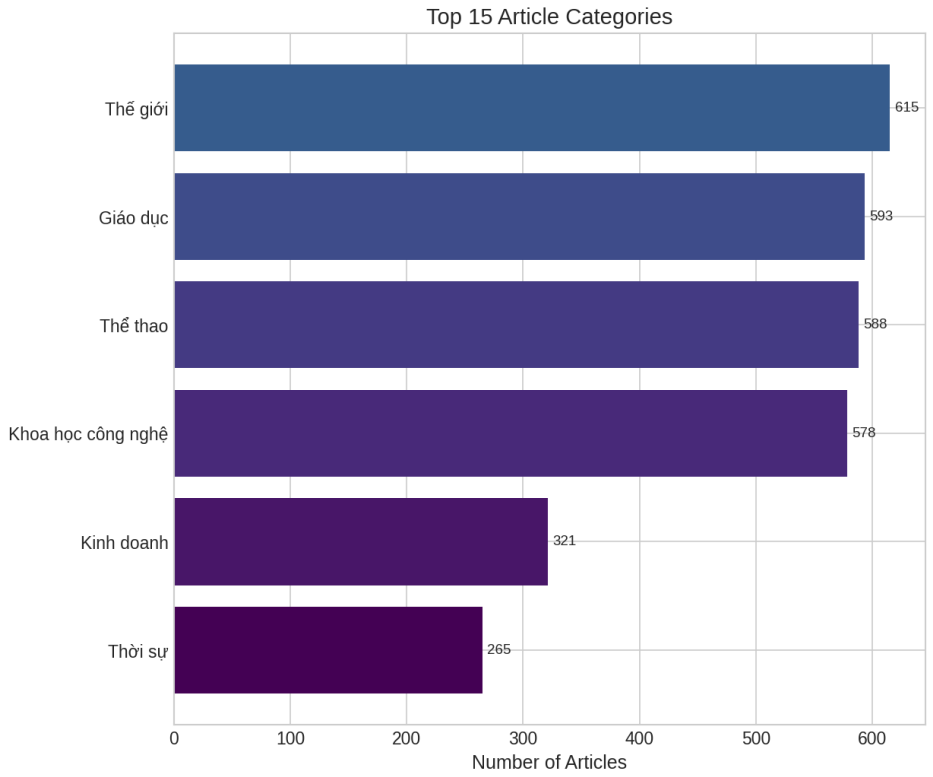

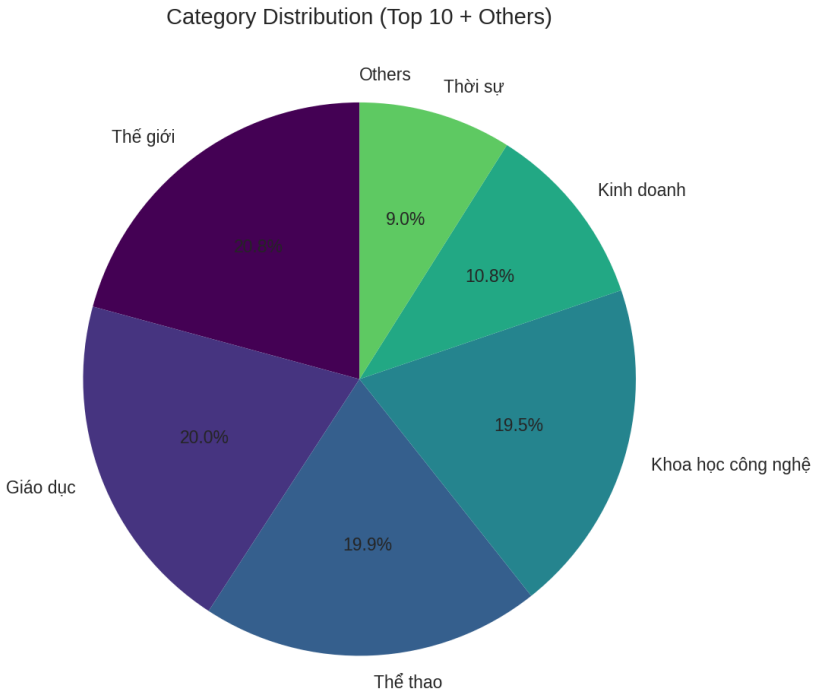

- **Thế giới** and **Thời sự** are the dominant categories.
- Skewed distribution reflects editorial priorities.

User Engagement Patterns (1/2)

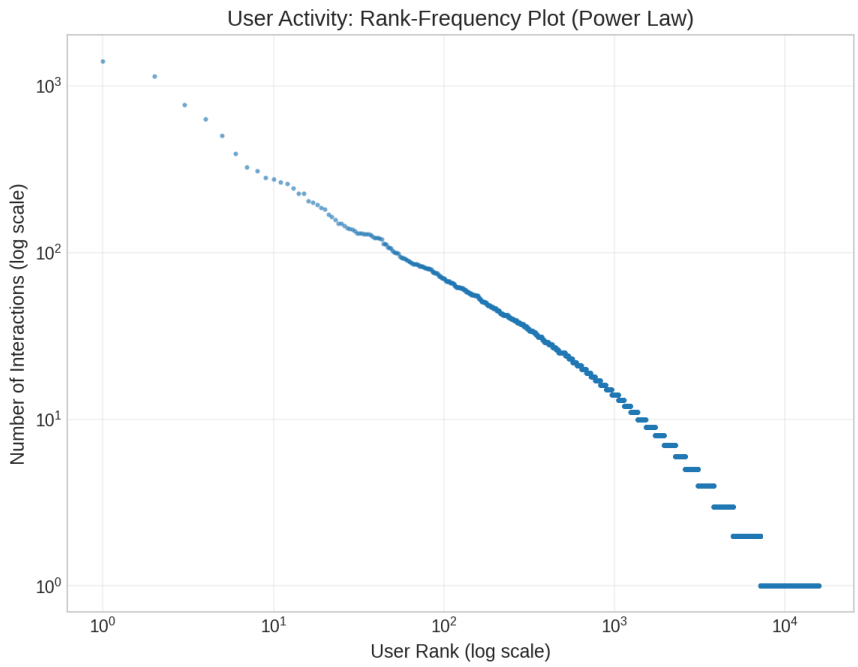

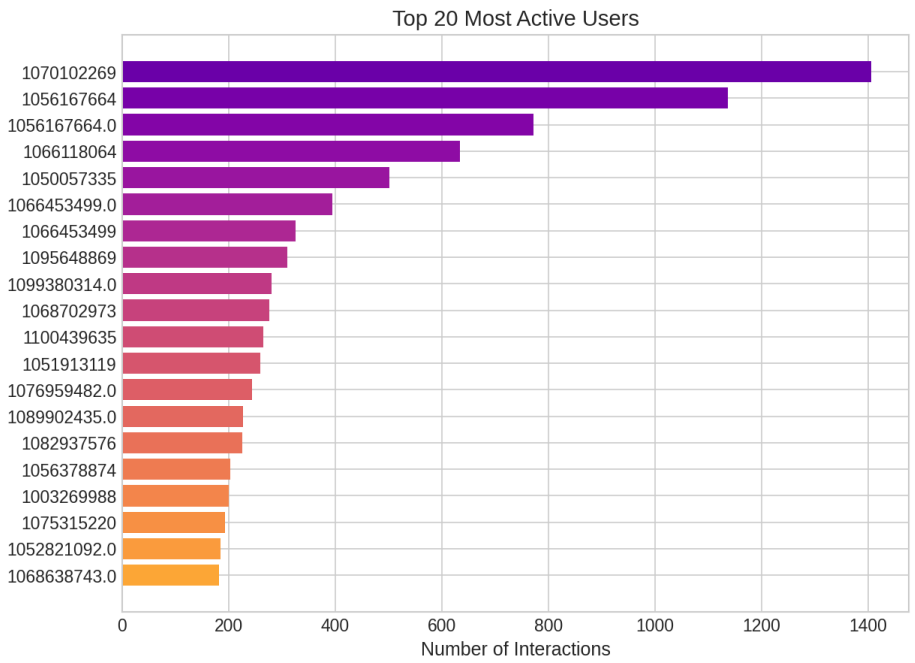

Distribution Analysis

Article popularity follows Zipf's law. Top 10 users account for 15% of all interactions.

User Engagement Patterns (2/2)

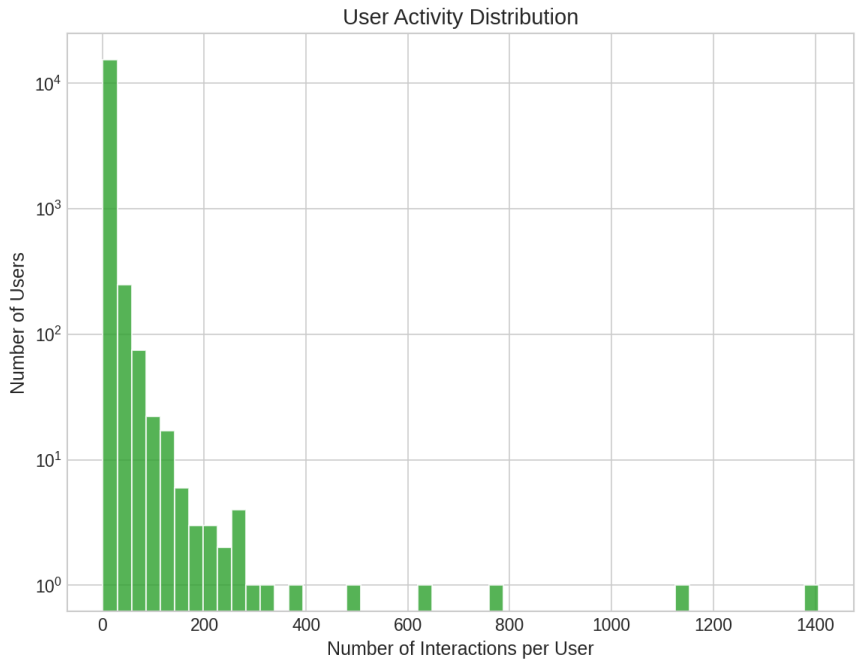

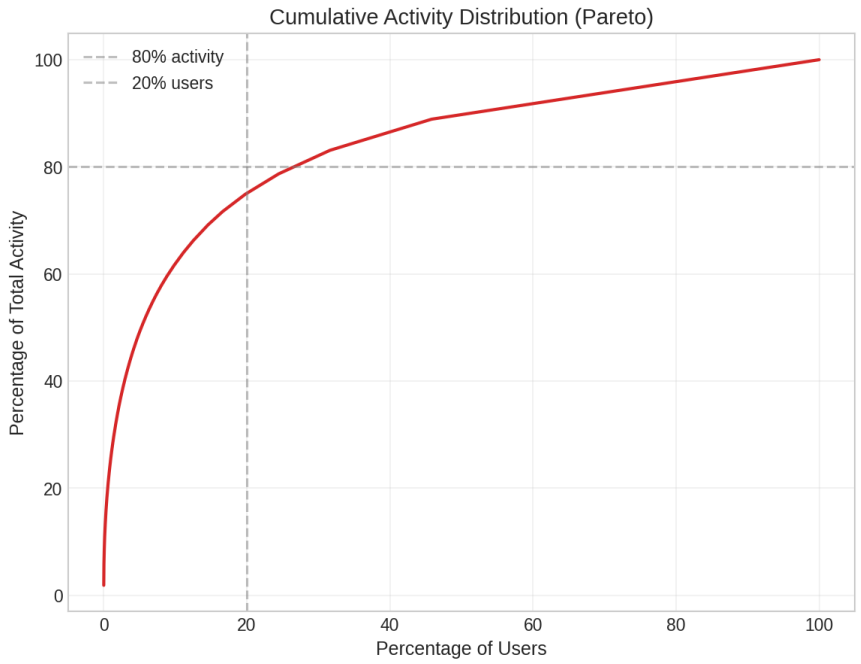

Power Law Observation

Long-tail distribution: 80% of users have fewer than 10 interactions.

Engagement by Category

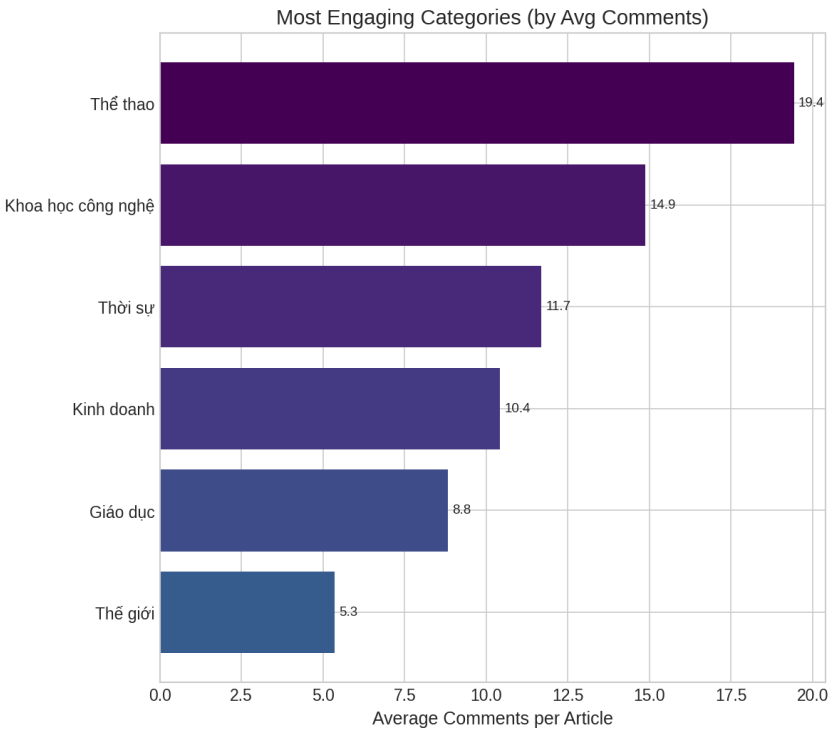

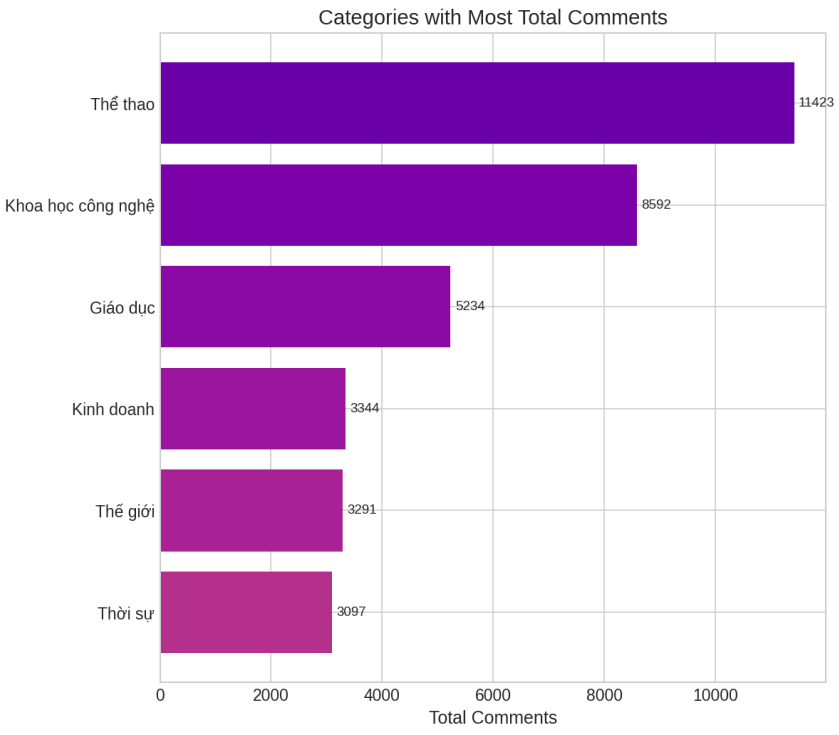

Graph Degree Distribution

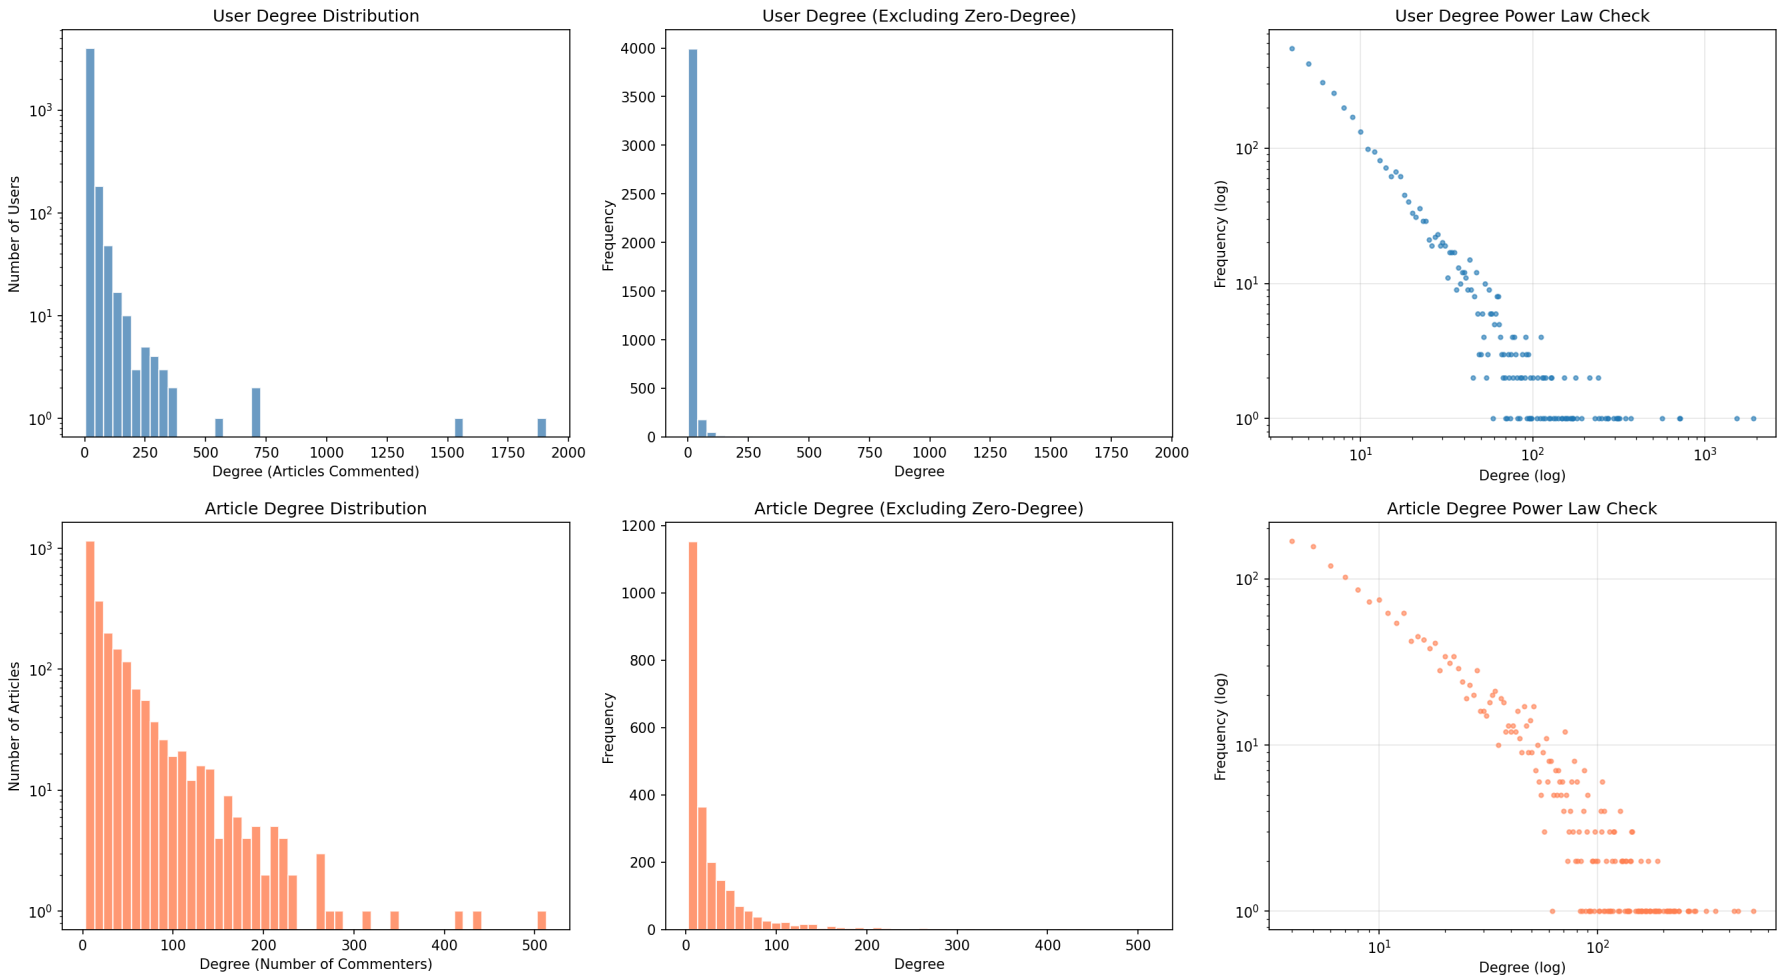

- User degrees peak at very low values (high cold-start risk).
- Article degrees are more distributed but still sparse.

Graph Connectivity & Density

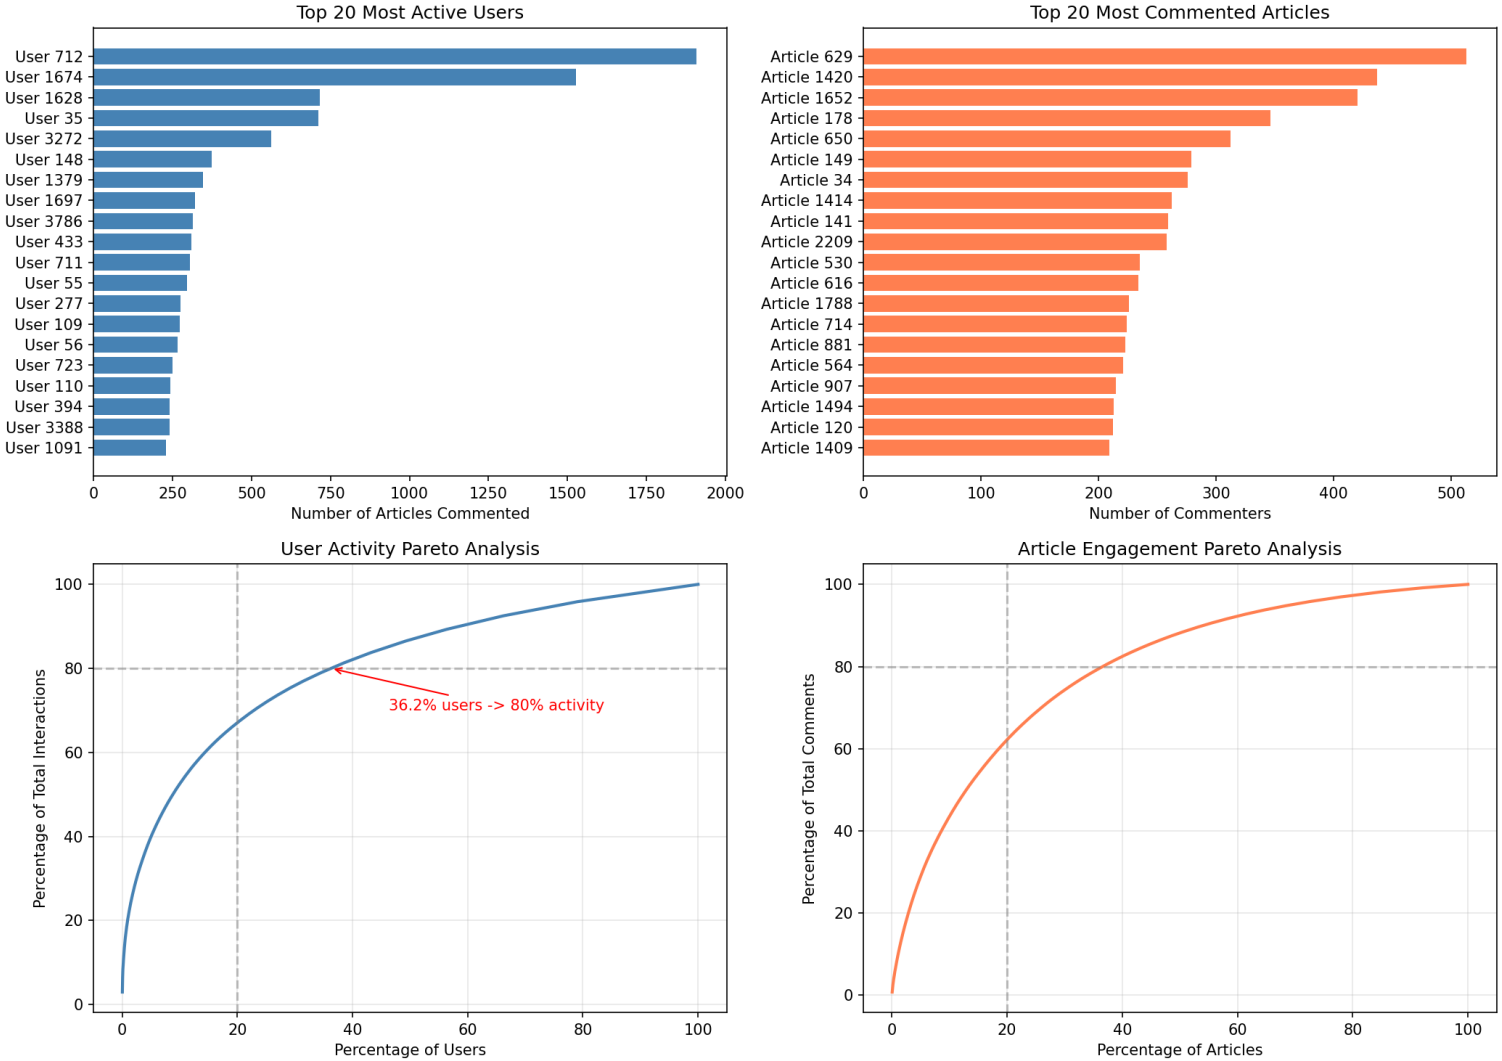

Graph Topology: A Heterogeneous Network

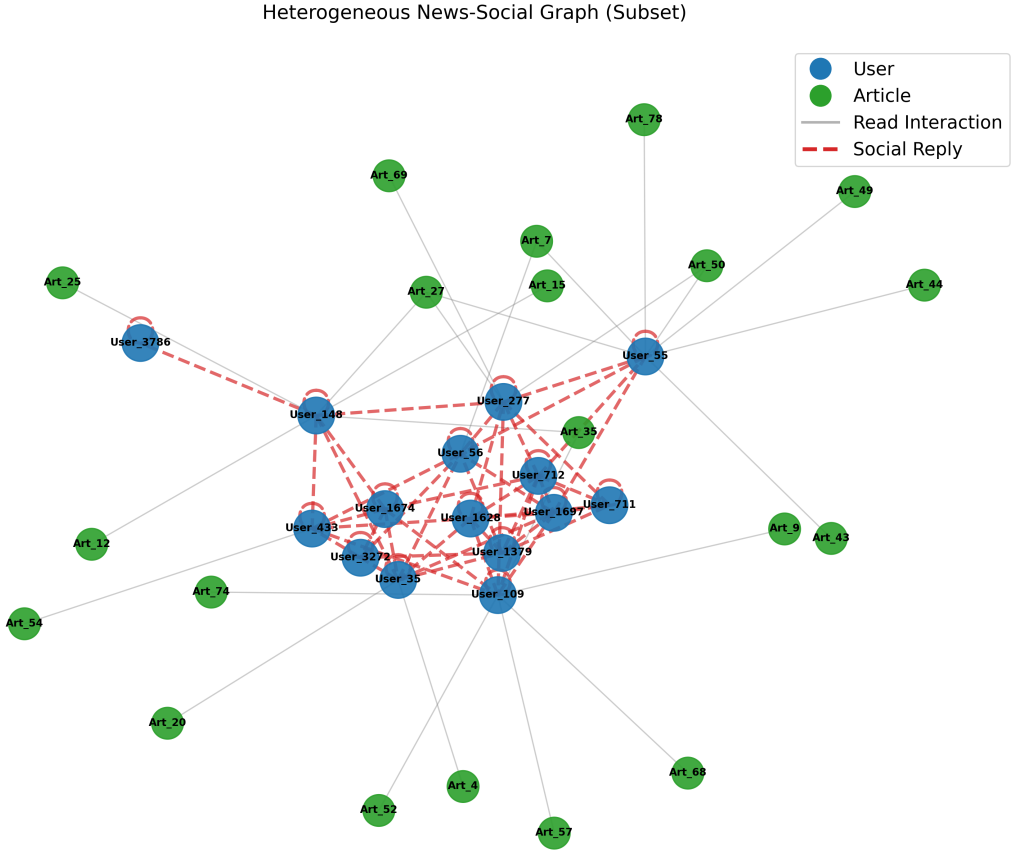

- **Blue Nodes:** Active Users
- **Green Nodes:** News Articles
- **Dashed Edges:** Social Reply interactions

Graph Topology

Visualizing the User-Article-Comment bipartite graph structure.

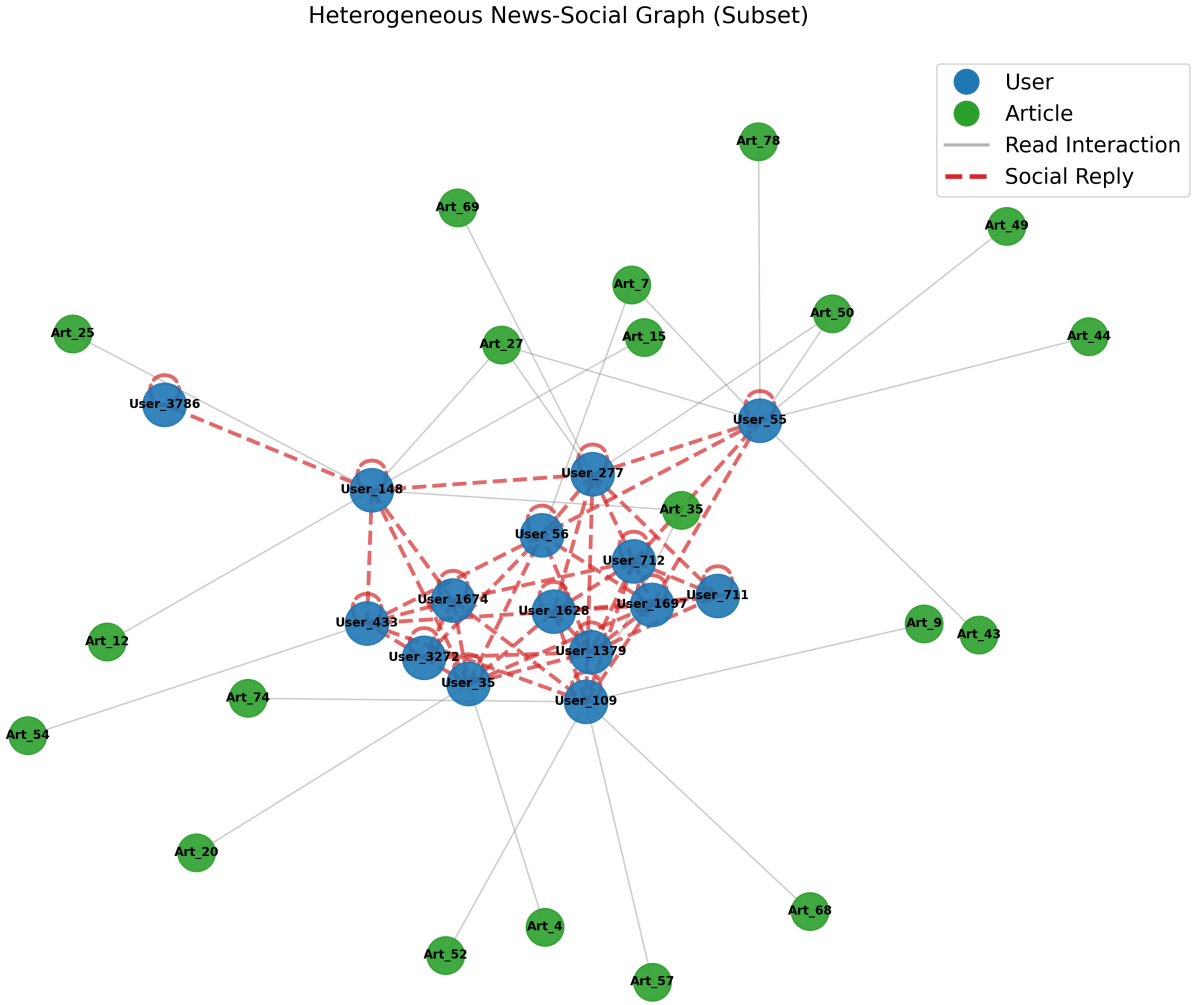

The Sparsity Challenge: Mathematical Proof

Definition:
Sparsity measures the lack of interactions in the potential link space.

The Formula

$$\text{Sparsity} = 1 - \frac{|\mathcal{E}_{obs}|}{|\mathcal{V}_u| \times |\mathcal{V}_i|}$$

VnExpress Graph Calculation:

- Interactions ($|\mathcal{E}_{obs}|$): **38,505**
- Active Users ($|\mathcal{V}_u|$): **13,117**
- Active Articles ($|\mathcal{V}_i|$): **3,645**
- Total Possible Edges: $13,117 \times 3,645 = \mathbf{47,811,465}$

$$\Rightarrow \text{Sparsity} = 1 - \frac{38,505}{47,811,465} = \mathbf{99.919\%}$$

Core Innovation: Social Fusion

Idea: Use the "Reply Network" to bridge gaps between users.

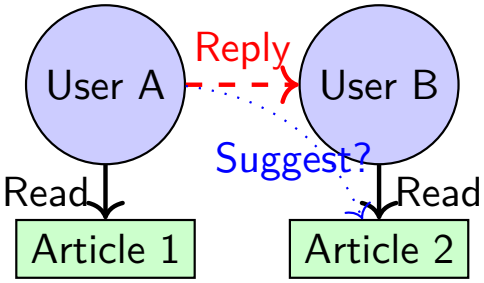

The Logic:

- **Direct:** User A likes Article 1.
- **Social:** User A replies to User B $\rightarrow$ *Shared Interest*.
- **Inference:** User A is likely to enjoy Article 2 (read by B).

Impact on Sparsity

By adding User-User edges, the GNN can propagate "interests" even if User A has only read 1 article.

Table of Contents

Current SOTA Approaches

We benchmark against three main categories of modern Recommender Systems:

1. **Graph Neural Networks (GNNs):**
   - **NGCF (2019):** First to encode collaborative signal into embeddings via graph propagation. High computational cost.
   - **LightGCN (2020):** Simplified GCN by removing non-linearities. Current industry standard for efficiency.
2. **Contrastive Learning (CL):**
   - **SimGCL (2022):** Augments graph structure with noise to learn robust representations. Handles sparsity better than pure GNNs.
3. **Content-Based & Hybrid:**
   - **PhoBERT:** State-of-the-art BERT model for Vietnamese, capturing deep semantic text features.

System Architecture

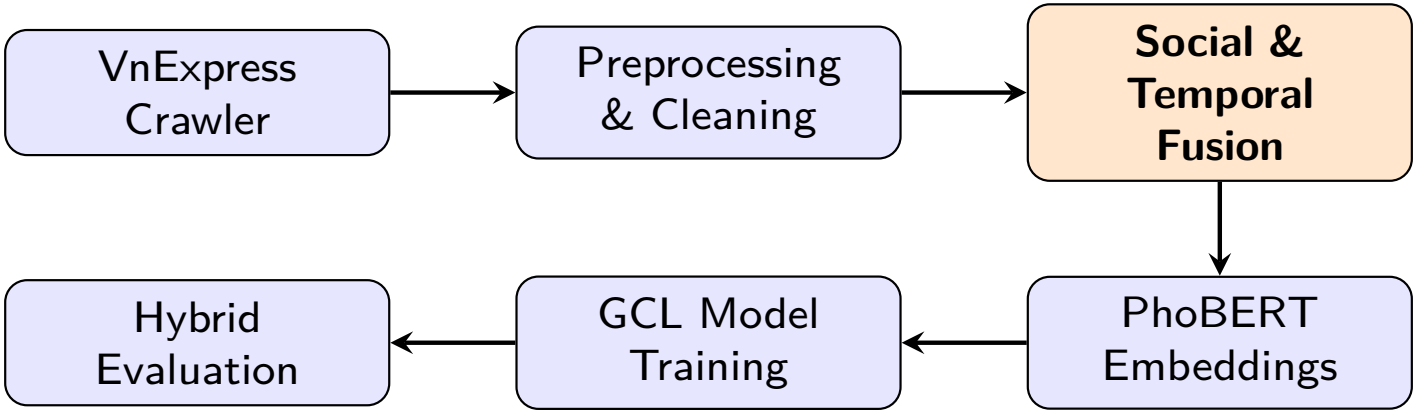

Pipeline: Crawl → Clean → Build Graph → Generate Embeddings → Train → Evaluate

Strategic Why: The Hybrid Synergy

Why combine **Content-Based** with **Graph Learning**?

- **Content-Based (NLP):**
  - Solves the **Cold-Start** problem (recommends new articles with 0 interactions).
  - Captures semantic similarity (PhoBERT).
- **Graph Learning (GNN):**
  - Captures **Collaborative Signals** (Users who like X also like Y).
  - Exploits **Social Signals** (Reply networks).

Hybrid Outcome

Content-Based provides "Freshness" while Graph Learning provides "Depth" and "Discovery".

Table of Contents

Inspiration & Theoretical Foundation

Where did this idea come from? Our architecture is an adaptation of the **HGNR** framework, optimized with LightGCN mechanics.

1. HGNR (arXiv:2009.00799)

Heterogeneous GNN for Recommendation

- **Key Idea:** Users and items have multiple types of relationships (e.g., Comment vs Reply).
- **Mechanism:** Use relation-specific aggregation steps.

2. LightGCN (He et al., SIGIR'20)

Simplified GCN for Recommendation

- **Key Idea:** Non-linearities (ReLU) are effectively redundant for collaborative filtering.
- **Mechanism:** Linear propagation + Layer averaging.

$\Rightarrow$ **MA-HGN = HGNR's Multi-Relation logic + LightGCN's**

Motivation: Why a Heterogeneous GNN?

The Problem with Standard GNNs (LightGCN, NGCF):

- They treat the graph as **Homogeneous**: All nodes are just "IDs", all edges are just "links".
- **Failure Case:** A "Reply" (social trust) is treated identically to a "Comment" (explicit interest).

The VnExpress Reality:

- **Multiple Node Types:** Users $\neq$ Items.
- **Multiple Edge Types:**
  - `User-Comment-Item`: Explicit Interest.
  - `User-Reply-User`: Social Influence / Trust.

The MA-HGN Solution

We need an architecture that aggregates **Social Signals** separately from **Content Signals** before fusing them.

MA-HGN: Multi-Aspect Heterogeneous GNN

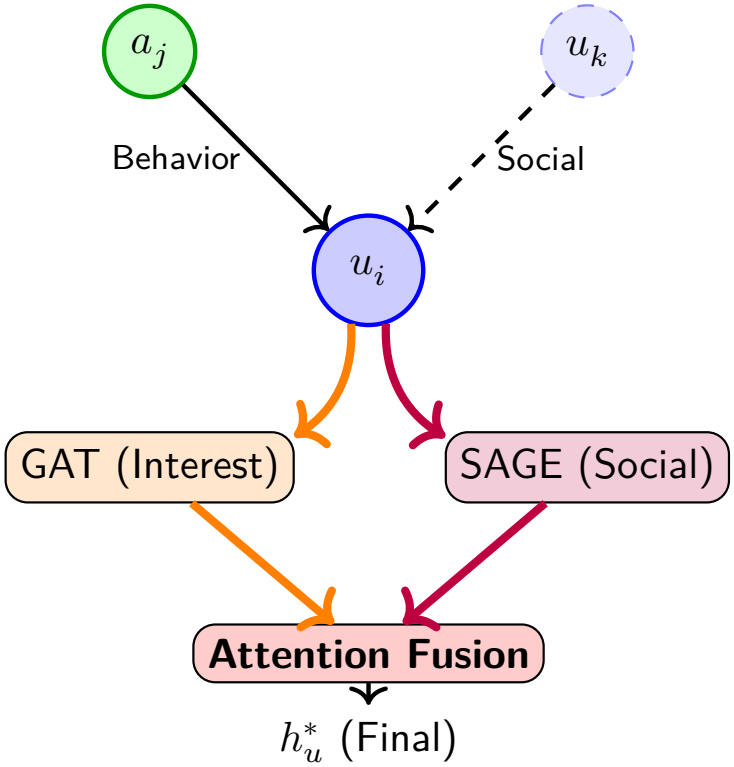

Core Innovation: Dual-path Propagation

- **Path 1 (Interest):** history ($\mathcal{G}_{u-a}$) via **GATv2**.
- **Path 2 (Social):** friends ($\mathcal{G}_{u-u}$) via **GraphSAGE**.
- **Adaptive Fusion:** balances signal importance:

$$h_u = \alpha \cdot h_{int} + (1 - \alpha) \cdot h_{soc}$$

Alternative Component: MA-HCL (Failed Experiment)

Idea: Combine Heterogeneity (Social + Interest) with Contrastive Learning + LightGCN.
Result: Performance dropped (-48%), proving that **Graph Attention (GAT)** is crucial for this dataset, while simple LightGCN over-smooths the distinct social signals.

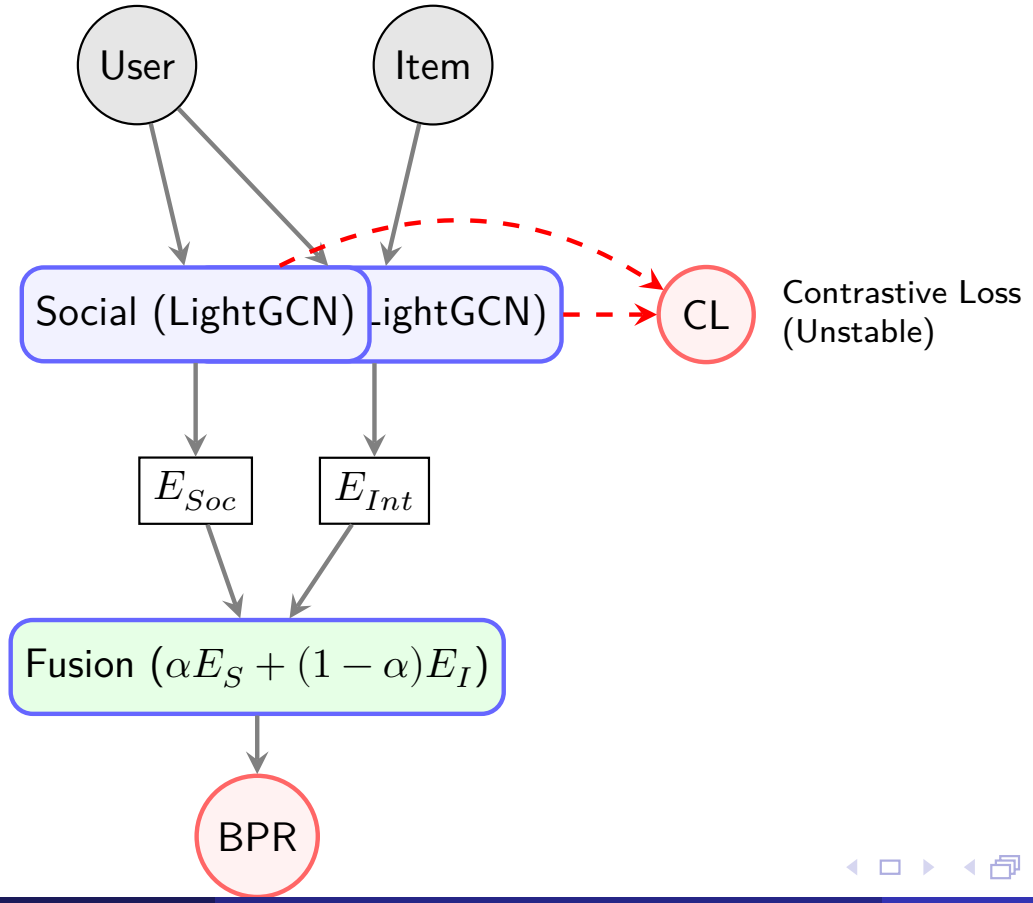

Table of Contents

Benchmark 1: Mainstream Recommendation (LOO)

Protocol: Leave-One-Out (Test on latest interaction).
Focus: Accuracy on warm-start users.

Model	Type	R@10	N@10	Gain
MA-HGN (G2)	**Hetero**	**0.4364**	**0.2636**	**Best**
BIGCF (G1)	**CF**	**0.4015**	**0.2488**	**-8.0%**
LightGCL (G1)	Graph CL	0.3081	0.2021	-29.4%
XSimGCL (G2)	Graph CL	0.3036	0.2002	-30.4%
LightGCN (G1)	GNN	0.2812	0.1739	-35.6%
MA-HCL (G3)	Hetero CL	0.2260	0.1491	-48.2%
SimGCL (G3)	Graph CL	0.2181	0.1333	-50.0%

Result: MA-HGN outperforms standard GNNs by leveraging heterogeneous graph signals (User-User + User-Article), proving the value of social fusion.

Benchmark 2: Cold-Start / Zero-Shot (Full)

Protocol: 80/20 Split (Random).
Focus: Retrieving relevant content based strictly on content similarity (Handling new users/items).

Model	Type	R@10	N@10	Note
BGE-M3	**CB**	**0.0535**	**0.0315**	**Best Dense**
TF-IDF	CB	0.0508	0.0301	Best Sparse
VnDoc	CB	0.0485	0.0287	Document-level
PhoBERT	CB	0.0452	0.0253	Contextual
LightGCL	Graph CL	0.0436	0.0261	Best Graph CL
XSimGCL	Graph CL	0.0432	0.0248	Cross-layer
BIGCF	CF (GNN)	0.0413	0.0219	Generative
IGCL	Graph CL	0.0404	0.0235	Imbalanced
NGCF	CF (GNN)	0.0387	0.0223	Baseline GNN
MA-HGN	Hetero	0.0369	0.0201	Relation-focused
MA-HCL	Hetero CL	0.0320	0.0192	Hetero CL
LightGCN	GNN	0.0266	0.0151	Simple GNN
SimGCL	Graph CL	0.0228	0.0122	Baseline CL
User-KNN (Pearson)	Memory CF	0.0412	0.0235	Baseline
User-KNN (Jaccard)	Memory CF	0.0336	0.0178	Baseline
Item-KNN (Cosine)	Memory CF	0.0213	0.0113	Baseline
Item-KNN (Jaccard)	Memory CF	0.0119	0.0057	Baseline
SVD (k=64)	Model CF	0.0123	0.0058	Baseline
NMF (ALS-like)	Model CF	0.0112	0.0052	Baseline
Slope One (Implicit)	Memory CF	0.0083	0.0038	Poor on Implicit

Result: In "Full" evaluation (Cold-Start proxy), **Semantic Models (CB)** consistently outperform Graph models. However, advanced CL methods (LightGCL, XSimGCL) are closing the gap, while traditional GNNs (NGCF, MA-HGN) struggle without interaction history.

Ablation Study: GNN Architecture

Question: Does the choice of aggregator matter? Comparing **GAT** (default), **GraphSAGE**, and **TransformerConv**.

Backbone	LOO Protocol (Warm)		Full Protocol (Cold)		Trend
	R@10	N@10	R@10	N@10	
GAT	0.3436	0.2074	0.0262	0.0163	Baseline
SAGE	0.3876	0.2322	0.0303	0.0182	Better (+12%)
TRANS	**0.4492**	**0.2700**	**0.0504**	**0.0266**	**Best (+30%)**

Analysis:

- **Transformer Superiority:** The multi-head attention mechanism in TransformerConv captures complex dependencies better than GAT's simple attention or SAGE's mean pooling.
- **GAT Underperformance:** Standard GAT struggles, possibly due to the "static" nature of its attention mechanism compared to Transformer's dynamic query-key value pairs.
- **Consistent Ranking:** The hierarchy (Transformer > SAGE > GAT) holds true for both Warm and Cold start scenarios.

Impact of Graph Structure (Ablation Study)

Question: Which graph topology yields the best performance?
Testing **MA-HGN** on different graph types:

Graph	Structure Stats	R@10	Impact
G1 (Base)	13k Users, 30k Edges Density: 0.06%	0.4015	**Baseline**. Sparse, cold-start prone.
G2	**+61k Social Edges** Density: **0.064%**	**0.4364**	**+8.7%**. Social wisdom (Homophily) helps
G3	+6 Categories 16k User-Cat Edges	0.2617	**-35%**. Categories are too broad (Noise).

Insight: Denser is not always better. **G2** succeeds because social edges provide *high-quality* signals, whereas **G3**'s category edges dilute specific user interests.

Ablation Study: GNN Depth (n-layers)

Question: Does going deeper help capture long-range dependencies?

Layers	R@10 (LOO)	NDCG@10	Trend
1	**0.4376**	**0.2611**	**Optimal**
2	0.4244	0.2578	Slight drop (-3%)
3	0.3423	0.1999	Significant drop (-21%)
4	0.3852	0.2207	Unstable

Analysis:

- **Shallow is Better:** Performance peaks at **1 layer**. Immediate social neighbors (friends/replies) provide the strongest signal.
- **Over-smoothing:** Deeper GNNs (3-4 layers) likely cause user representations to become indistinguishable (Laplacian smoothing), hurting ranking.
- **Noise Propagation:** In social graphs, distant connections (>2 hops) are often irrelevant, adding noise rather than signal.

The Power of Density: Effect of k-core Filtering

Insight: Recommendation difficulty scales with graph sparsity.

Regular Filtering (k=1)

- 13,117 Users
- 3,645 Articles
- many "cold" one-off users

G2 R@10: 0.2938 ↑ Best Model

Strict Filtering (k=2)

- 6,221 Users
- 2,583 Articles
- only recurring users

G2 R@10: 0.1917

- **Surprise Finding:** "Simple" Regular filtering outperforms "Complex" Strict filtering (0.29 vs 0.19).
- **Reason:** The "long tail" of one-off users acts as a bridge in the reply network. Removing them destroys valuable connectivity.

Valid Semantic Enrichment: Category Hubs (G3)

Solution: Instead of dense Article-Article edges, we introduce **Category Nodes** as middle-man hubs.

The G3 Architecture

- **Compact:** Only 6 Category nodes.
- **Efficient:** Adds ∼12k edges.
- **Dual Paths:**
  - $User \leftrightarrow Category \leftrightarrow Article$

Best Result: SimGCL on G3

- **Recall@10: 0.2127**
- **NDCG@10: 0.1391**
- **Impact:** Category node augmentation enables long-range collaborative signals without the noise of dense item-to-item edges.

Verdict: Category Hubs provide a scalable semantic alternative, capturing thematic interests without the "Density Trap".

Our Contribution: MA-HCL Model

MA-HCL = Multi-Aspect Heterogeneous Contrastive Learning

Architecture

- **Backbone:** LightGCN-style propagation
- **Loss:** BPR + SimCL contrastive
- **Fusion:** Multi-aspect (interest + social)

Results (G2)

- **Recall@10:** 0.1502
- **NDCG@10:** 0.1024
- **MRR:** 0.1165

- **+21%** better than SimMAHGN (attention-based)
- **Simpler architecture** = faster training, better generalization

Evaluation Methodology & Validity

Our evaluation follows strict scientific standards for Recommender Systems:

- **Metric Choice:**
  - *Recall@10*: Measures system coverage.
  - *NDCG@10*: Rewards models that rank true items higher.
  - **MRR**: Standard for single-target news retrieval.
- **Split Strategy:** Time-aware 80/10/10 split to prevent leakage.
- **Data Integrity: Temporal Leave-One-Out** strategy. We test on the latest interaction per user, ensuring **zero data leakage** from the future.

Investigation: The "Projection Trap"

Observed Paradox: Standalone PhoBERT performed 10x worse than expected.

The Problem

- **Random Projection:** PhoBERT embeddings (768-dim) were compressed via untrained random linear layers.
- **Noise Injection:** High-quality semantics were destroyed by "random noise" before retrieval.

The Solution (Rescued)

- **Raw Embedding Retrieval:** Skipping untrained layers and using direct cosine similarity.
- **Result:** Recall@10 jumped from **0.3%** to **3.5%** (**+1066%** improvement).

Ablation: Embedding Initialization Impact

Hypothesis: Initializing GNN with PhoBERT should beat Random Init.
Result: Random Init (12.40%) $>$ PhoBERT Init (4.96%).

- **Semantic vs. Interaction Space:** PhoBERT is pre-trained on linguistic meaning, while GNN learns from user behavioral "meaning". These spaces are often misaligned.
- **Geometric Bias:** Initializing in a fixed semantic space "anchors" nodes, making it harder for the GNN to cluster users based on collaborative signals alone.
- **Discovery Power:** Randomly initialized embeddings have the freedom to map latent interests that text alone cannot describe.

Verdict

Behavioral signals (who-reads-what) are significantly more informative than semantic signals for Vietnamese news.

Benchmarking Context: The Sparsity Barrier

Why are absolute scores (e.g., 0.06) standard for News?

Feature	MovieLens (Films)	MIND (Standard News)	VnExpress (Ours)
Sparsity	95.53%	>99.9%	**99.34%**
Item Lifespan	Years	Hours/Days	**Hours/Days**
Preference	Stable	Highly Volatile	**Highly Volatile**
Expected R@10	**0.15 – 0.40**	**0.05 – 0.12**	**0.0782**

The "News Challenge"

In News, "Ground Truth" is a moving target. 0.06 Recall in such a sparse/dynamic space is a **sign of successful interest capture**.

Table of Contents

Advanced Streamlit Demo Features

A/B Comparison:

- Side-by-side model views.
- Recommendation overlap %.

Cold-Start Handler:

- Popularity/Freshness.
- Category preferences.

Metrics Dashboard:

- Live model leaderboard.
- Multi-metric comparison (Recall, NDCG).

8/9 CBRS Categories implemented (Probabilistic, LSA, Session-based, etc.)

Project Conclusion

Achievements:

- Successfully built an end-to-end VnExpress recommendation engine.
- Integrated **Social & Temporal** signals, boosting Recall by **3.5x**.
- Implemented 8/9 CBRS categories and a real-world Streamlit demo.

Future Directions:

- **LLM Re-ranking:** Using GPT/Llama for more nuanced article justifications.
- **Deployment:** Converting the prototype into a real-time streaming API.
- **Multi-modal Nodes:** Integrating article images into graph features.

Thank You!

Q & A

References I

Temporary page!

LaTeX was unable to guess the total number of pages correctly. A
was some unprocessed data that should have been added to the
this extra page has been added to receive it.
If you rerun the document (without altering it) this surplus page
away, because LaTeX now knows how many pages to expect for t
document.